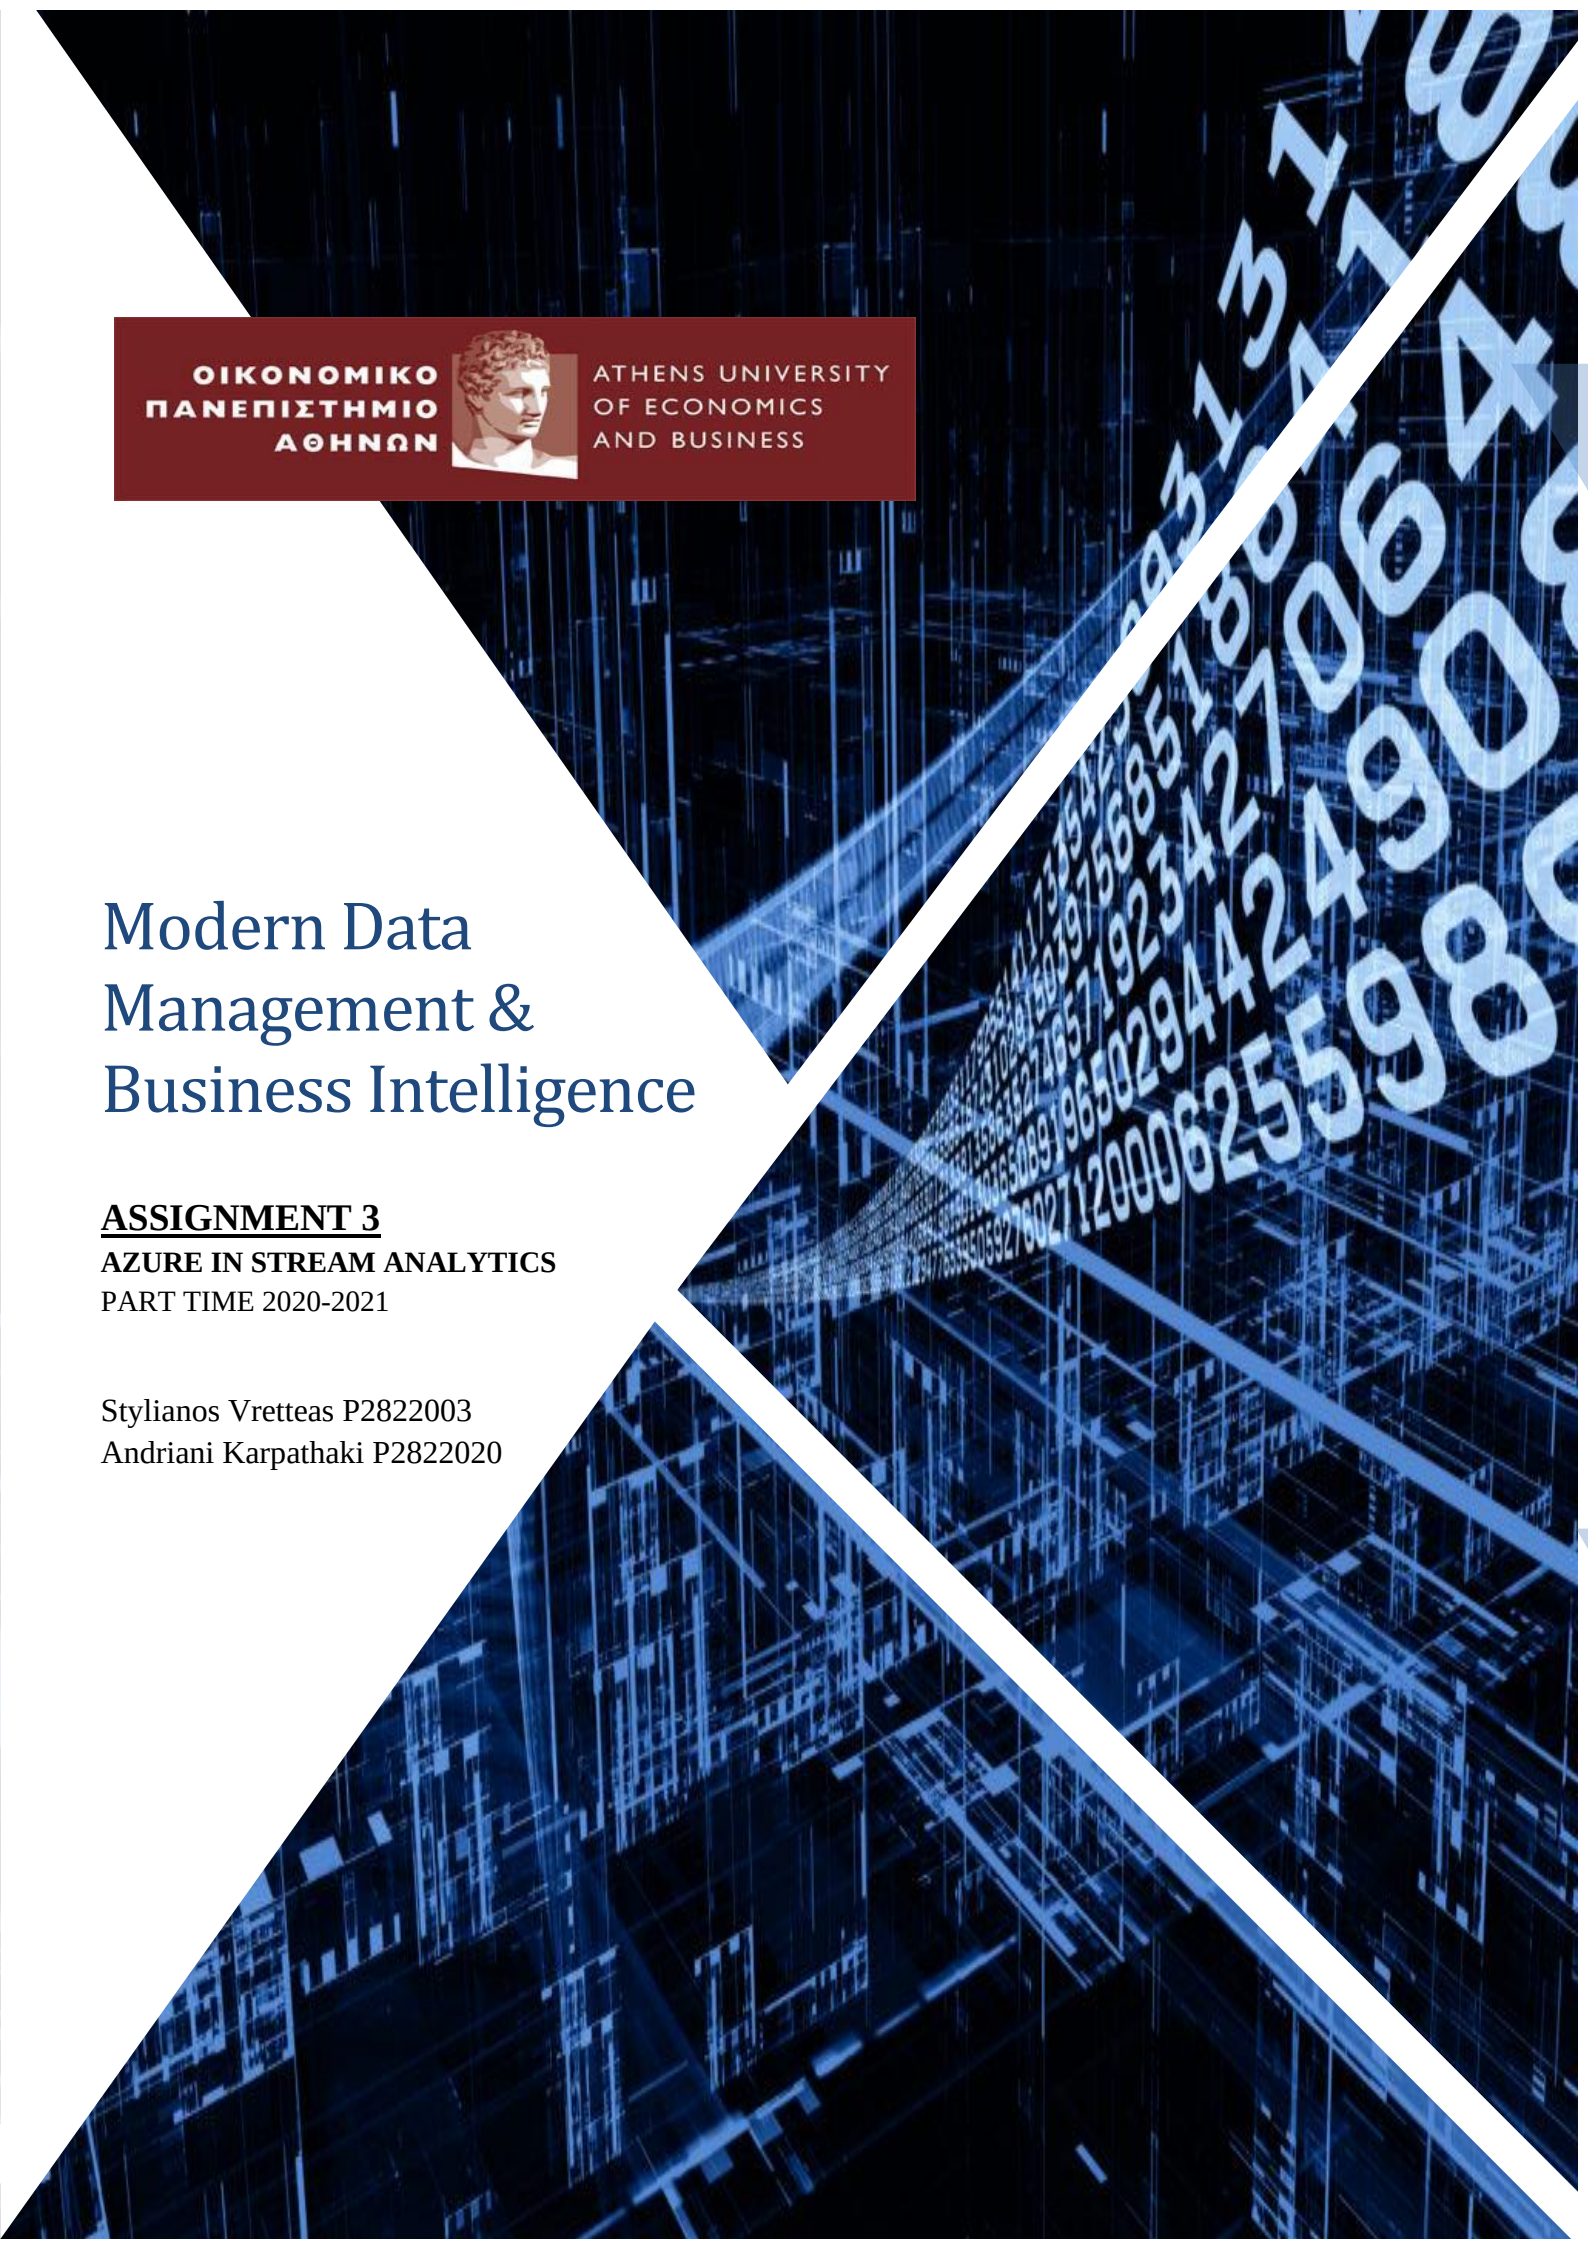

ΟΙΚΟΝΟΜΙΚΟ ΠΑΝΕΠΙΣΤΗΜΙΟ ΑΘΗΝΩΝ

ATHENS UNIVERSITY OF ECONOMICS AND BUSINESS

Modern Data Management & Business Intelligence

ASSIGNMENT 3

AZURE IN STREAM ANALYTICS

PART TIME 2020-2021

Stylianos Vretteas P2822003
Andriani Karpathaki P2822020

Contents

Abstract

The scope of this report was to use Azure Stream Analytics to process a data stream of ATM transactions and answer stream queries. The schema of the stream is: (ATMCode, CardNumber, Type, Amount). The steps followed in this procedure are represented below.

1. Create a student's account at: https://azure.microsoft.com/en-us/free/students/

First of all, we created an Azure Student account (Figure 1) and then we signed in (Figure 2).

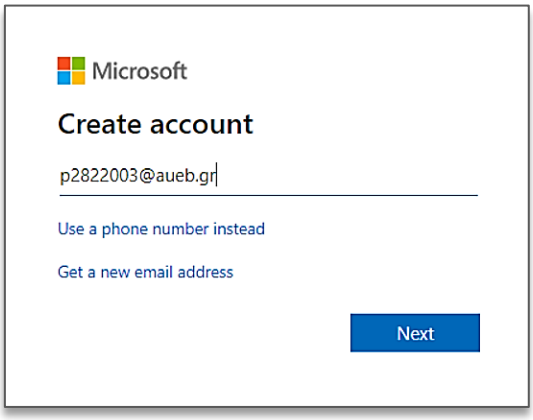

Figure 1 - Create an Azure Student Account

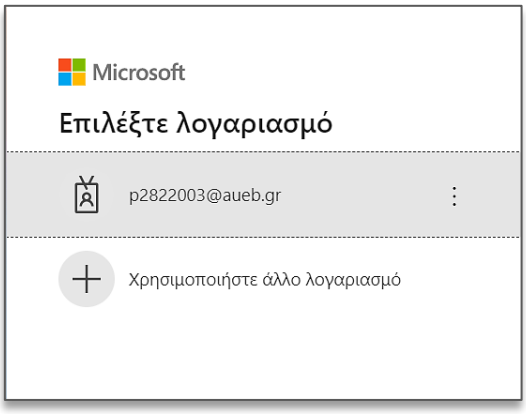

Figure 2 - Sign in Azure account

2. Setup an Event Hub.

We created a Resource group because it serves as a folder that holds related resources for an Azure solution. The resource group can include all the resources for the solution, or only those resources that are preferable to be managed as a group (Figure 3).

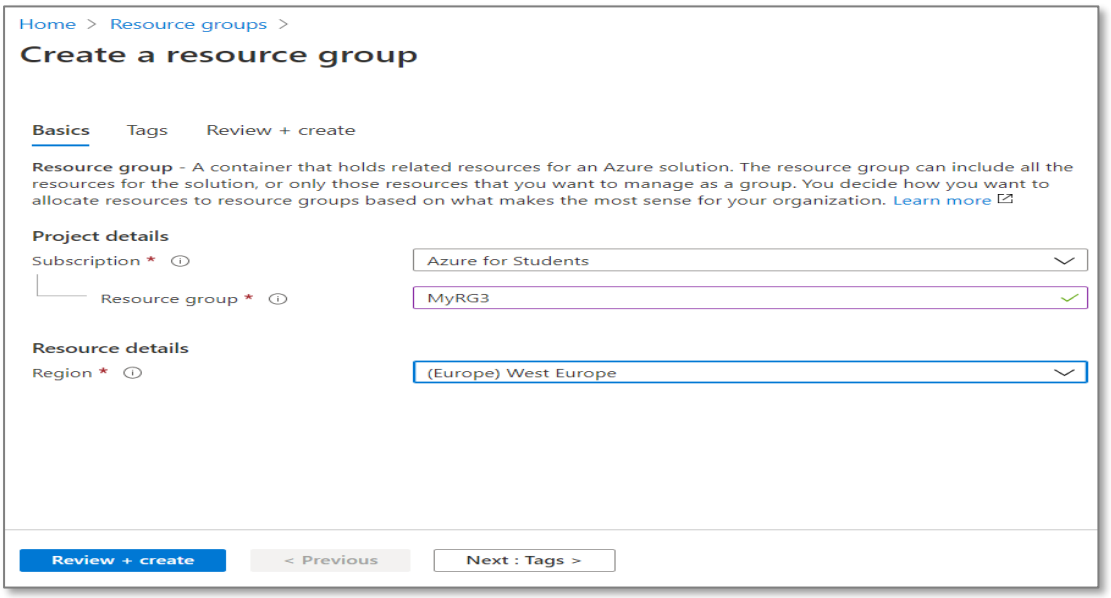

Figure 3 - Create a resource group

Next step was to create an Azure Event Hub which is a big data streaming platform and event ingestion service. It can receive and process millions of events per second. Data sent to an event hub can be transformed and stored by using any real-time analytics provider or batching/storage adapters (Figure 4).

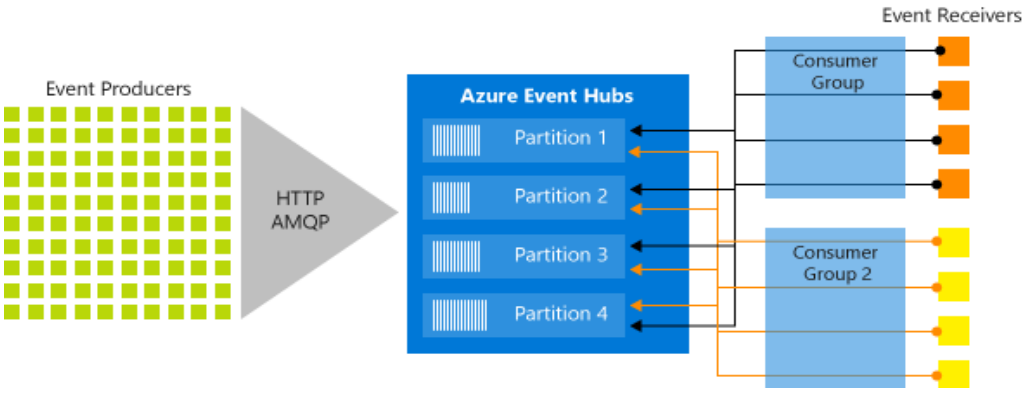

Figure 4 - Architecture Azure Event Hubs

Construction of Event Hub

At this point it necessary to create an Event Hubs Namespace. An Event Hubs namespace provides a unique scoping container, in which we can create one or more event hubs. For the scope of this paper, we have created one Event Hub, named MyAuebHub3 (Figure 5). The procedure was completed successfully (Figure 6).

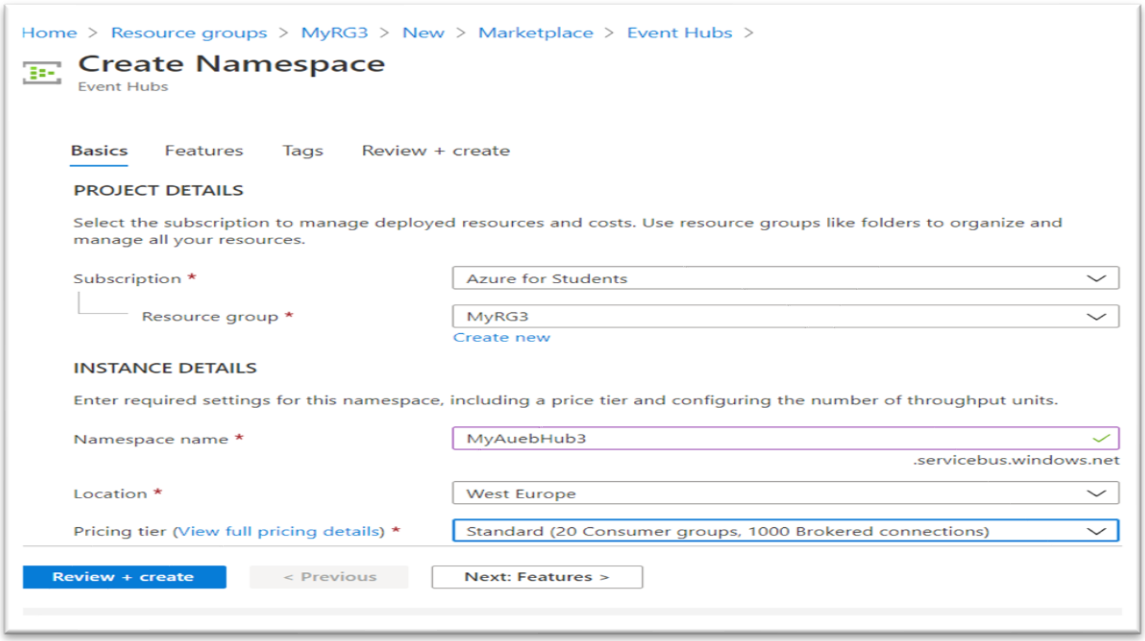

Figure 5 - Event Hub Creation

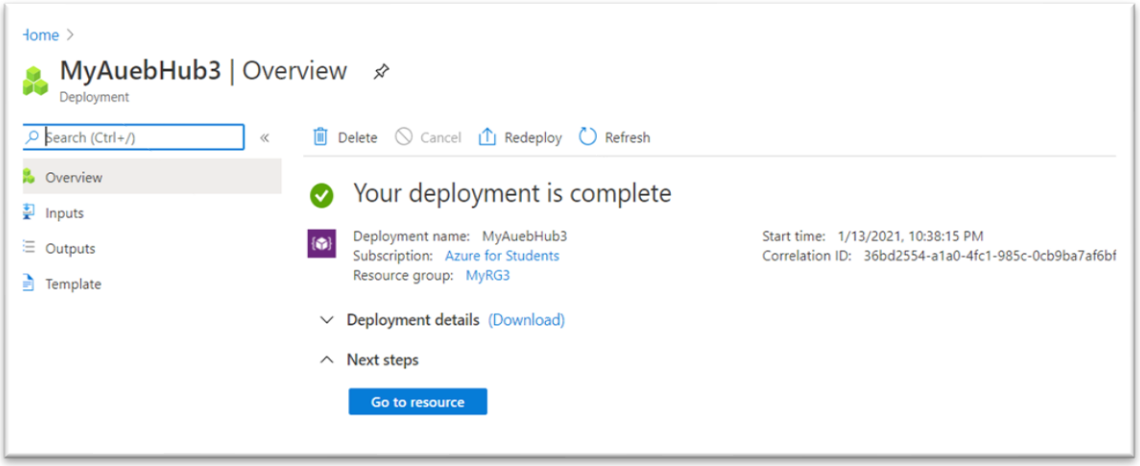

Figure 6 - Deployment completed for Event Hub creation

Next, we created an Event Hub instance (Figure 7).

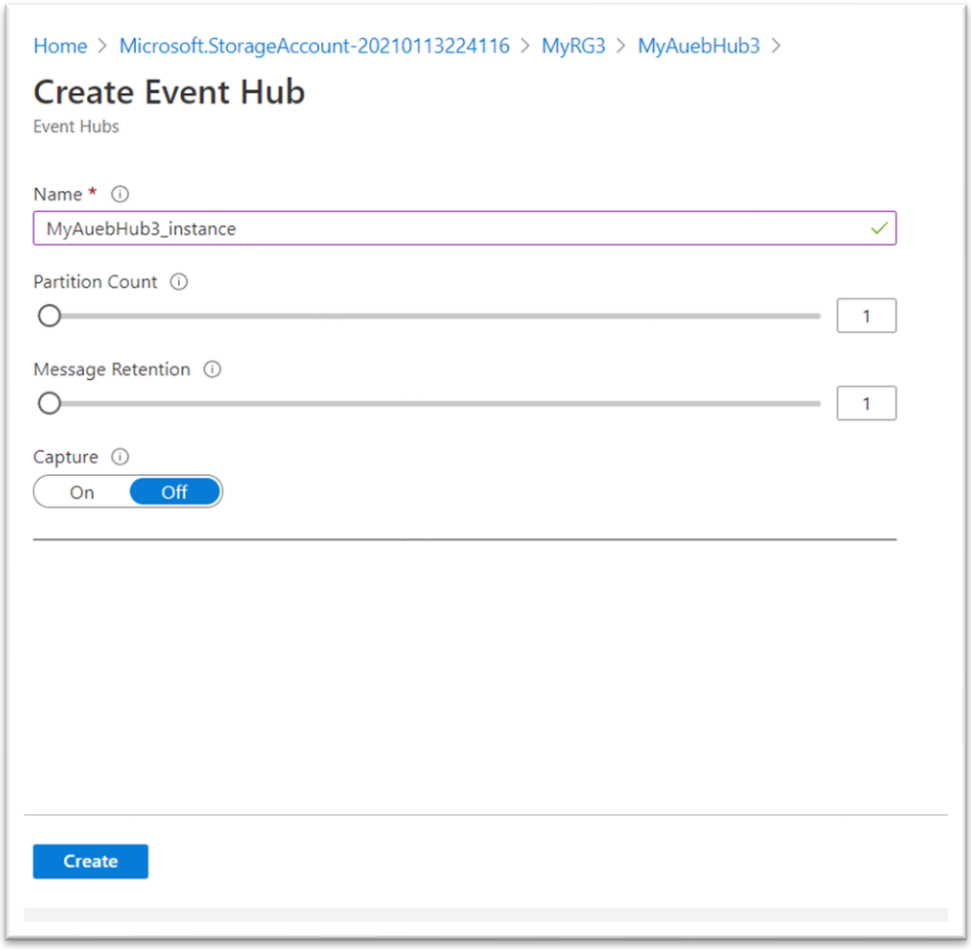

Figure 7 Event Hub Instance creation

3. Generate a Security Access Signature (use a terminal with windows operationg system): https://github.com/sandrinodimattia/RedDog/releases

Next step was to generate a security access signature and use it in order to edit the data transmission generator configuration settings and connect it with the event hub (Figure 8 & 9).

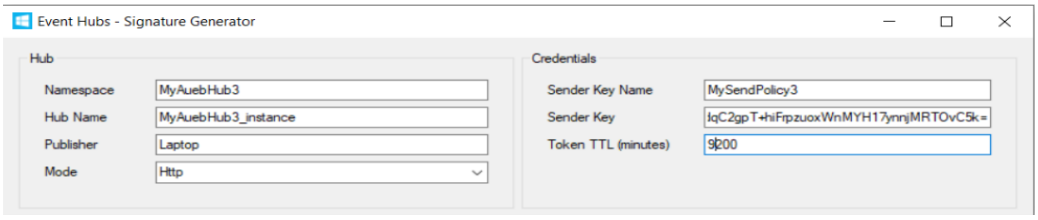

Figure 8 - Signature Generator

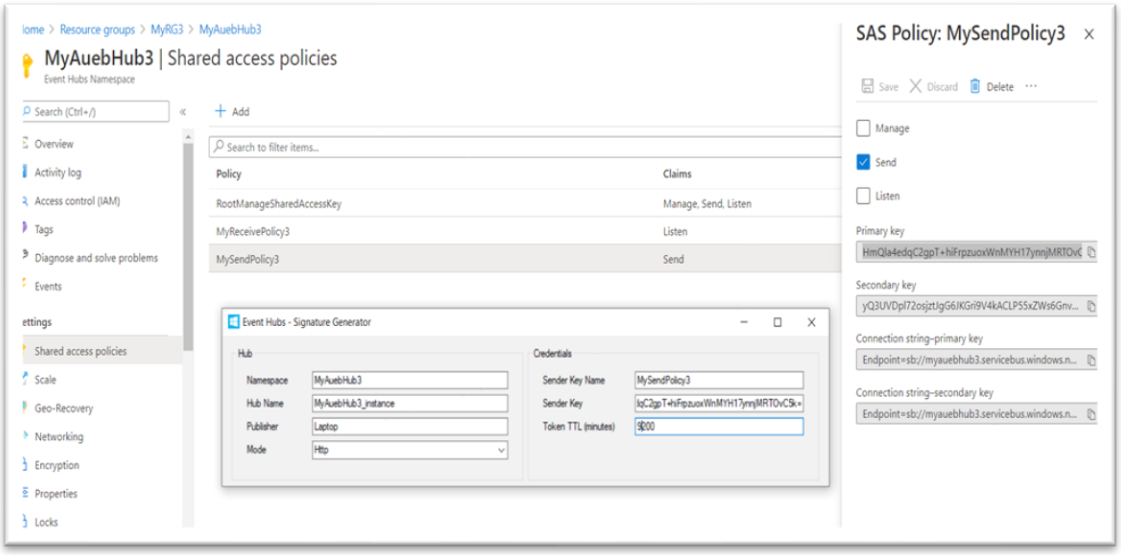

Figure 9 - Creation of signature

4. Edit Generator.html (open with a text editor, e.g.: Sublime or Notepad++) and update the CONFIG variables. Keep the "js" folder in the same folder as the Generator.html file.

After creating the signature, we edited the html file via Notepad++ and updated the CONFIG variables (Figure 10).

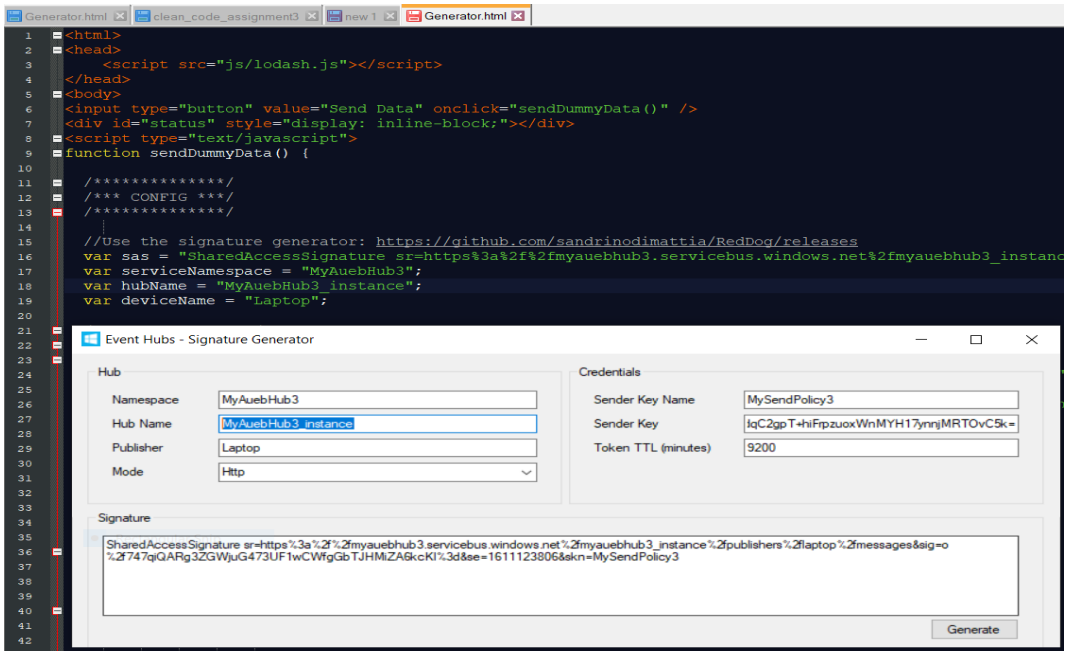

Figure 10 - Edit html file via Notepad++

5. Feed the Event Hub with the use of Generator.html (In order to start the Stream Generator, open the Generator.html with a web browser, e.g.: Chrome and press the “Send Data” button.)

After editing the html file, we opened the Generator.html via Chrome web browser in order to start the transmission of data (Figure 11).

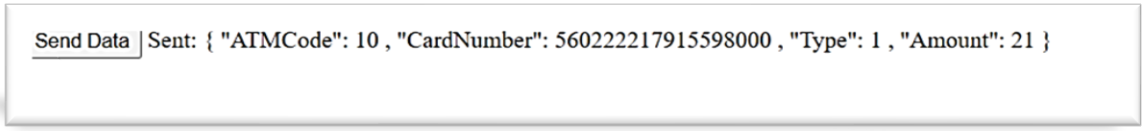

Figure 11 - Start of the Stream Generator

6. Setup a Storage account.

Storage Account

Azure Blob storage is Microsoft's object storage solution for the cloud. More precisely, Blob storage is optimized for storing massive amounts of unstructured data and offers three types of resources (Figure 12):

- The storage account which provides a unique namespace in Azure for your data
- A container in the storage account
- An unlimited number of containers as a blob in a container

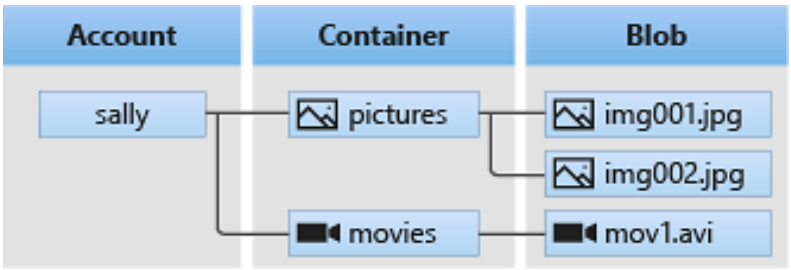

Figure 12 - Blob Storage Diagram

Construction of storage account and container

After creating the Event Hub, a storage account and container was built (Figure 13 & 15). To be noted that another container was created later in order to store the reference data there.

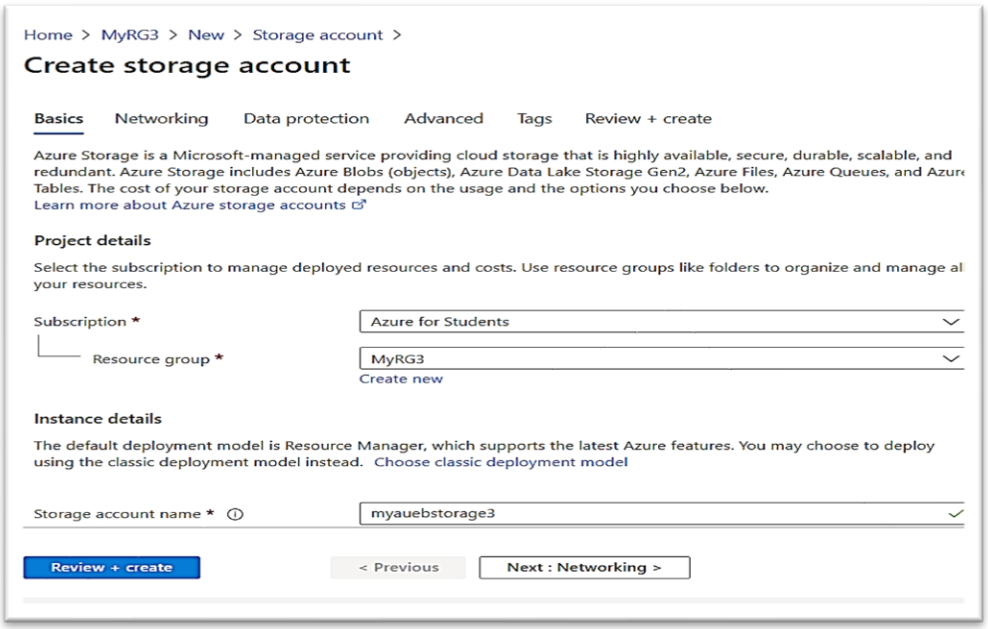

Figure 13 - Storage account creation

The procedure was completed successfully (Figure 14).

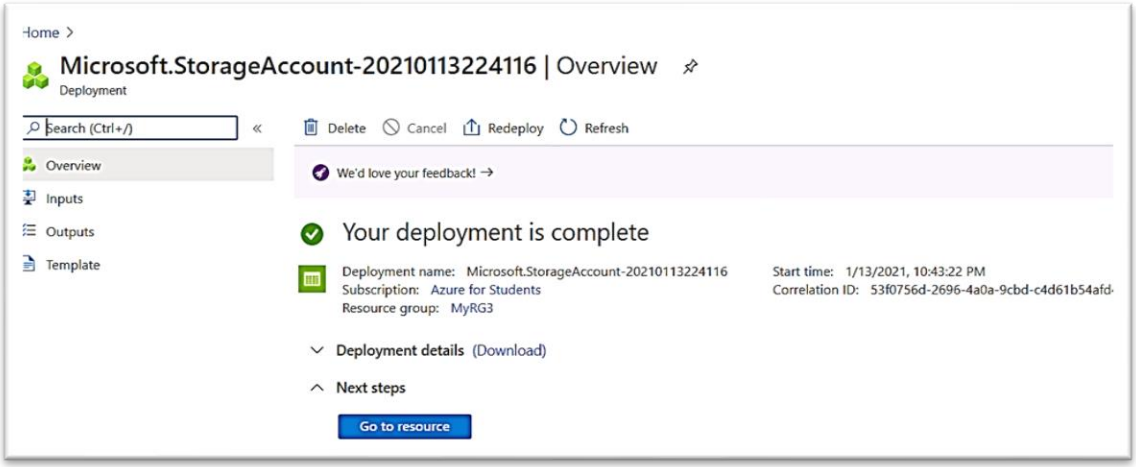

Figure 14 - Deployment completed for Storage Account creation

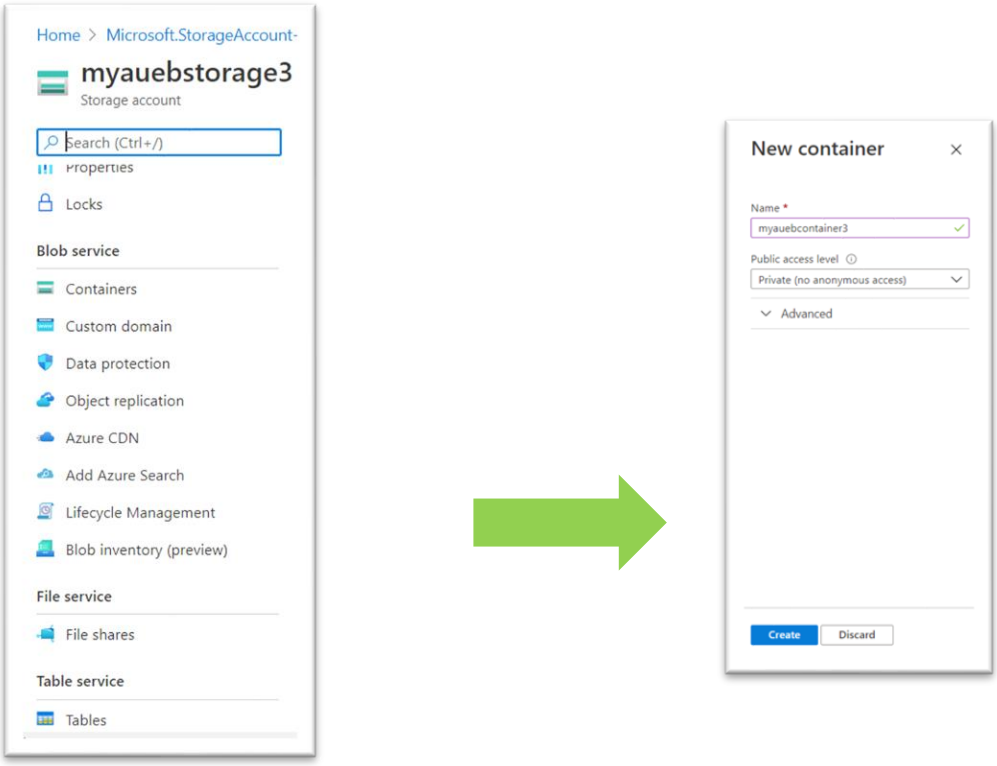

Figure 15 - Container creation

Construction of Policies

Afterwards, we constructed policies in the event hub so as to send and receive events. Thus , we created the MyReceivePolicy3 and the MySendPolicy3 policies (Figure 16).

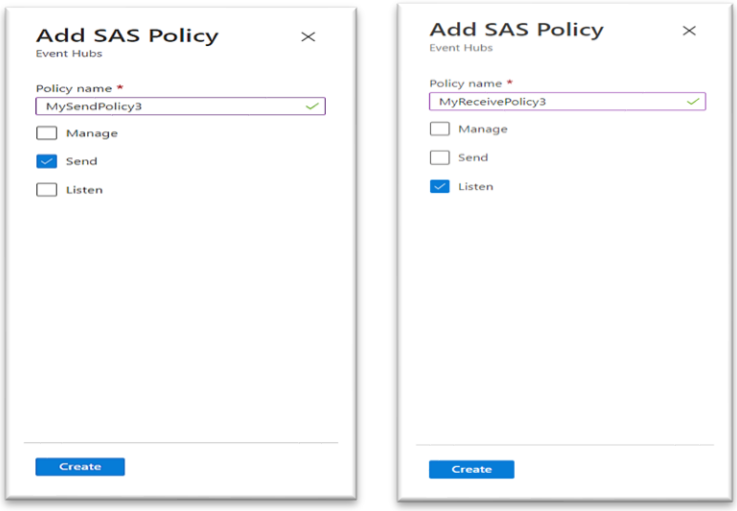

Figure 16 Creation of policies

7. Upload the Reference Data files to your storage account.

In order to implement queries with joins, we had to upload reference data files that would serve as lookup tables between the stream input and reference input. A second container was built with the purpose of storing the reference data (Figure 17 & 18).

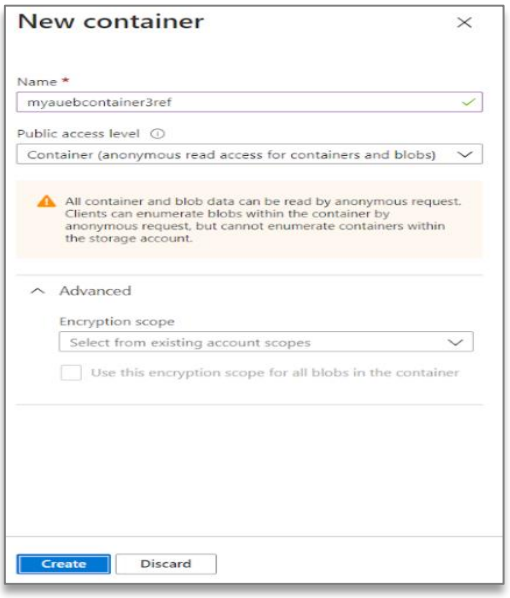

Figure 17 - Creation of a container for storing the reference data

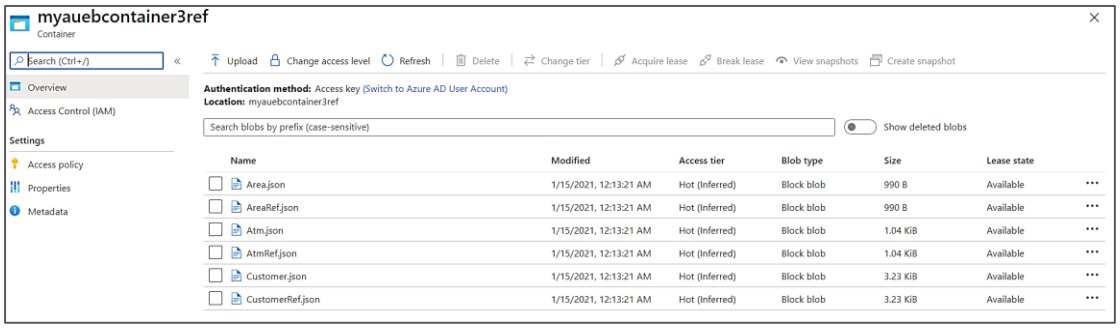

Figure 18 - Uploaded Reference Data

8. Setup a Stream Analytics Job.

Next, we had to create the last Azure solution which is the Stream Analytics Job (Figure 19 & 20). An Azure Stream Analytics job consists of an input, query, and an output. Below we can see the overview of the solutions created (Figure 21).

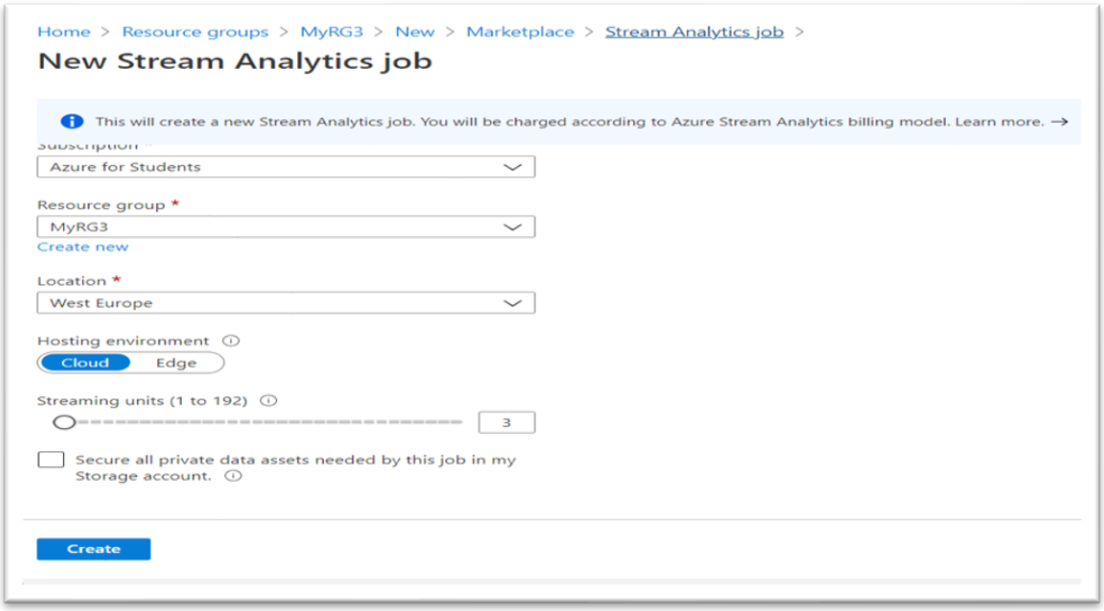

Figure 19 - New Streams Analytics job

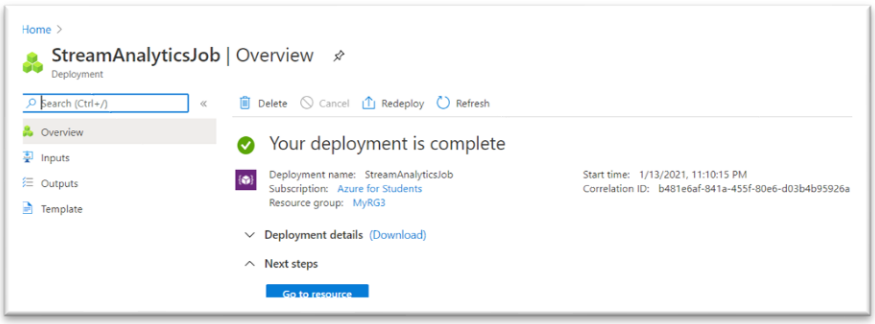

Figure 20 - Creation of New Streams Analytics job completed

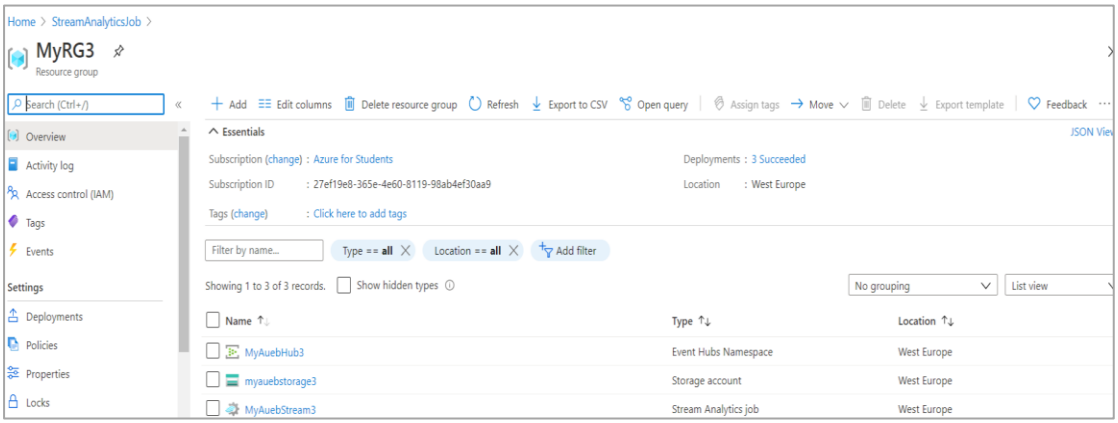

Figure 21 - Overview of the solutions created

9. Use the Event Hub + Reference Data Files as Input.

We had to define the input and output sources. There are two types of inputs:

Data stream input: A data stream is an unbounded sequence of events over time. Event Hubs are used to collect event streams from multiple devices and services.

Reference data input: Reference data is either completely static or changes slowly. It is typically used to perform correlation and lookups.

In our case for input we had one stream input regarding ATM transactions, three (3) reference inputs (Area, Atm, Customer) and one output (Figures 22-25).

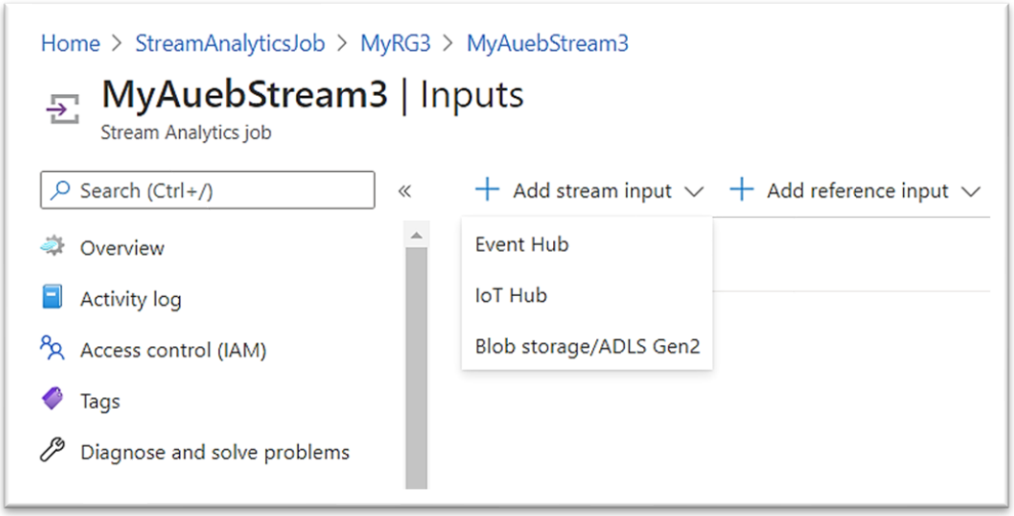

Figure 22 – Add stream input

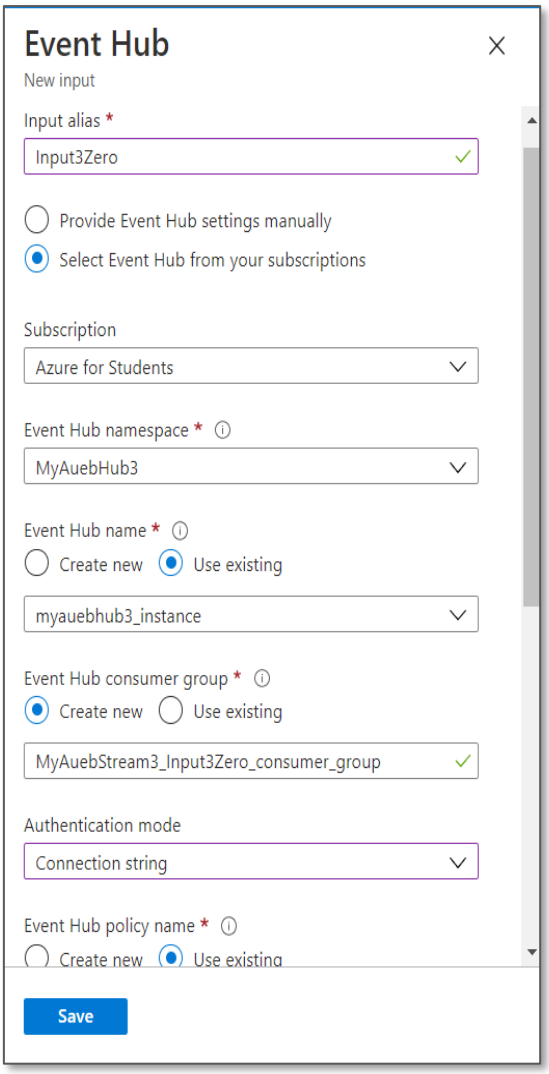

Figure 23 - Stream data input creation

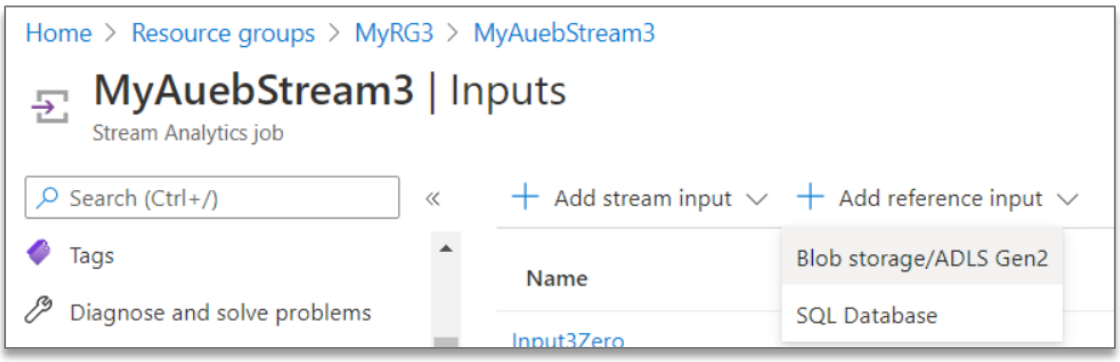

Figure 24 - Add reference data input

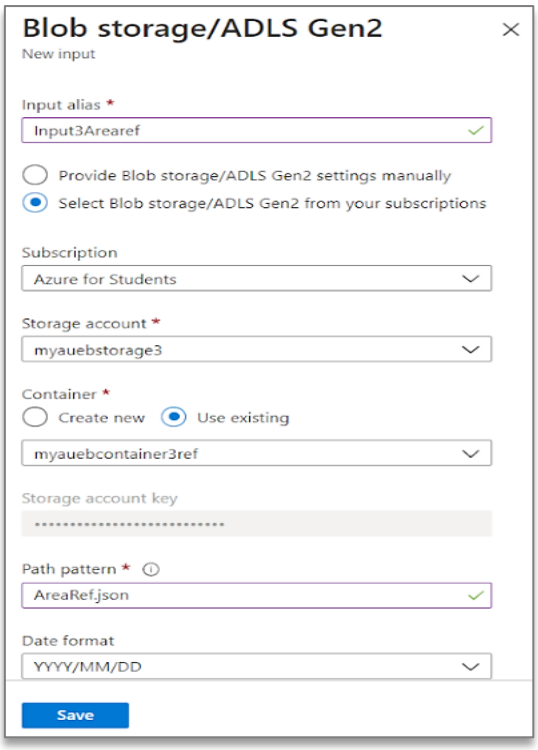

Figure 25 - Reference data input creation

10. Create a Blob Storage Output.

There are several output types for sending data. In our case we constructed a single output per job and linked it with the instance container (Figure 26).

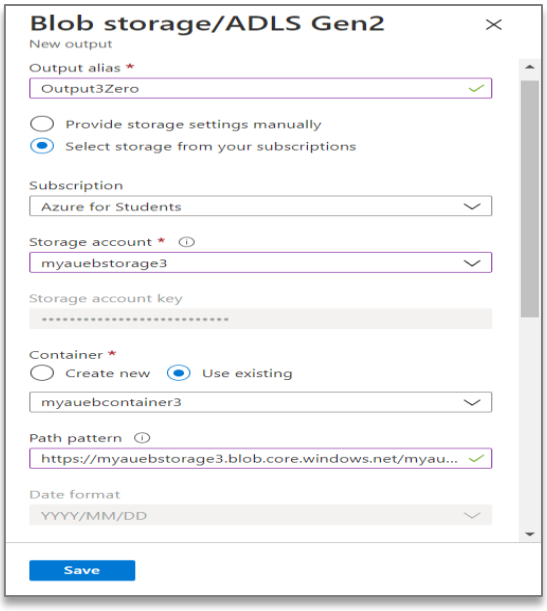

Figure 26 - Output creation

11. Queries

The SQL query language in Azure Stream Analytics gives the opportunity to carry out real-time analysis on streaming data. We created an Azure Analytics solution for the tasks listed in the "QUERIES" section, which comes as follows.

Query 1: Show the total "Amount" of "Type = 0" transactions at "ATM Code = 21" of the last 10 minutes. Repeat as new events keep flowing in (use a sliding window).

Query 1 - Code

```
select SUM(Amount) as Total_Amount
into [OutputZero]
from [InputZero]
timestamp by EventEnqueuedUtcTime
where Type = 0 AND ATMCode = 21
group by Type, ATMCode, SlidingWindow(minute, 10)
```

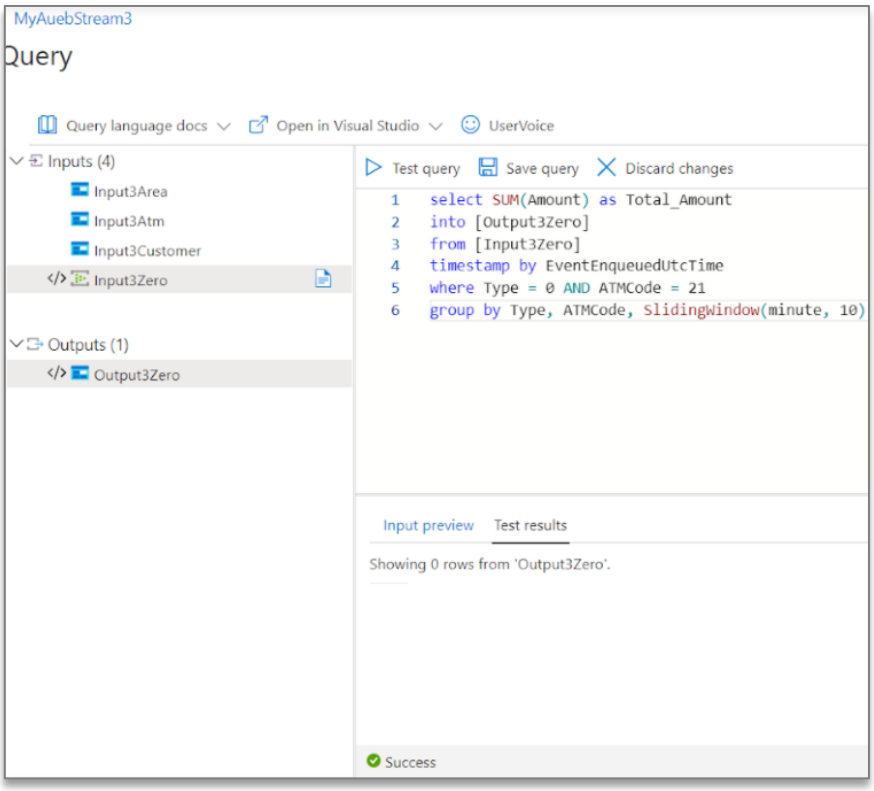

Figure 27 - Query 1 Test Results

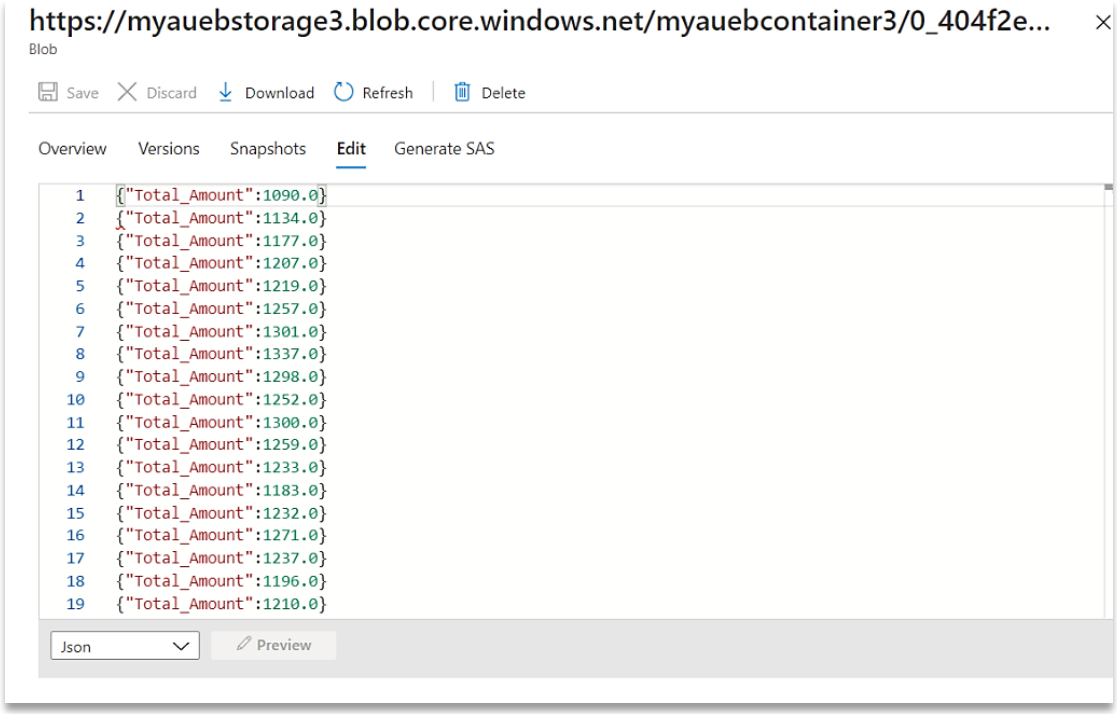

Figure 28 - Query 1 Output

Query 2: Show the total "Amount" of "Type = 1" transactions at "ATM Code = 21" of the last hour. Repeat once every hour (use a tumbling window).

Query 2 – Code

```
select SUM(Amount) as Total_Amount
into [Output3Zero]
from [Input3Zero]
timestamp by EventEnqueuedUtcTime
where Type = 1 AND ATMCode = 21
group by Type, ATMCode, TumblingWindow(hour, 1)
```

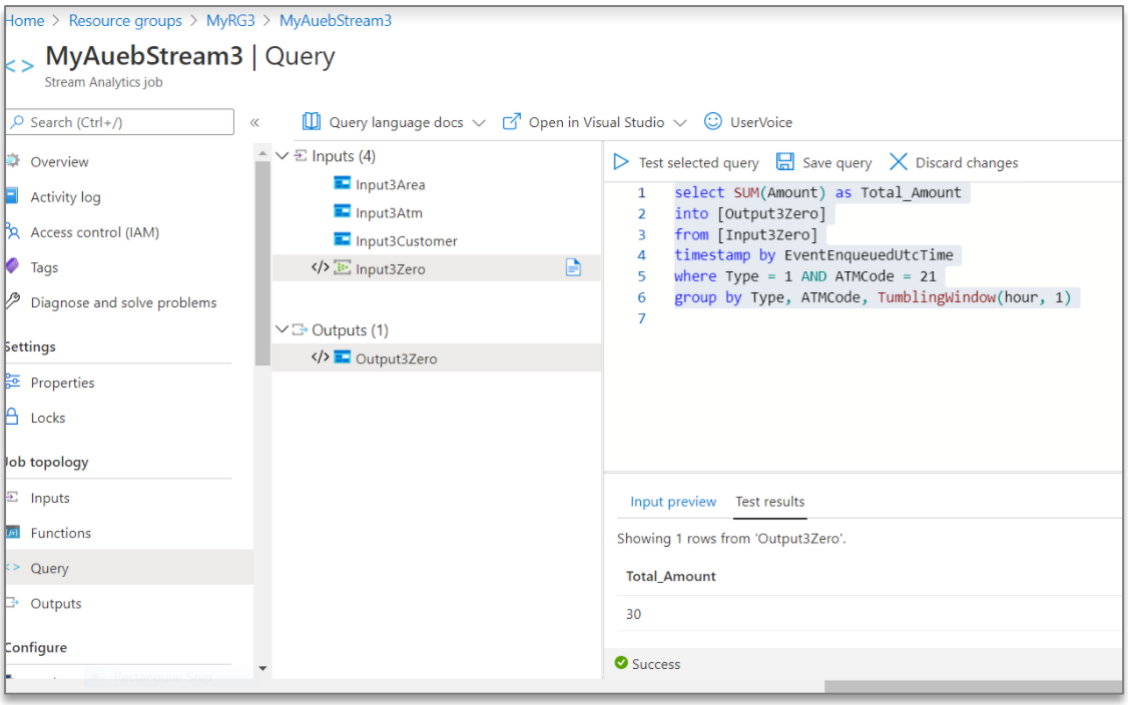

Figure 29 - Query 2 Test Results

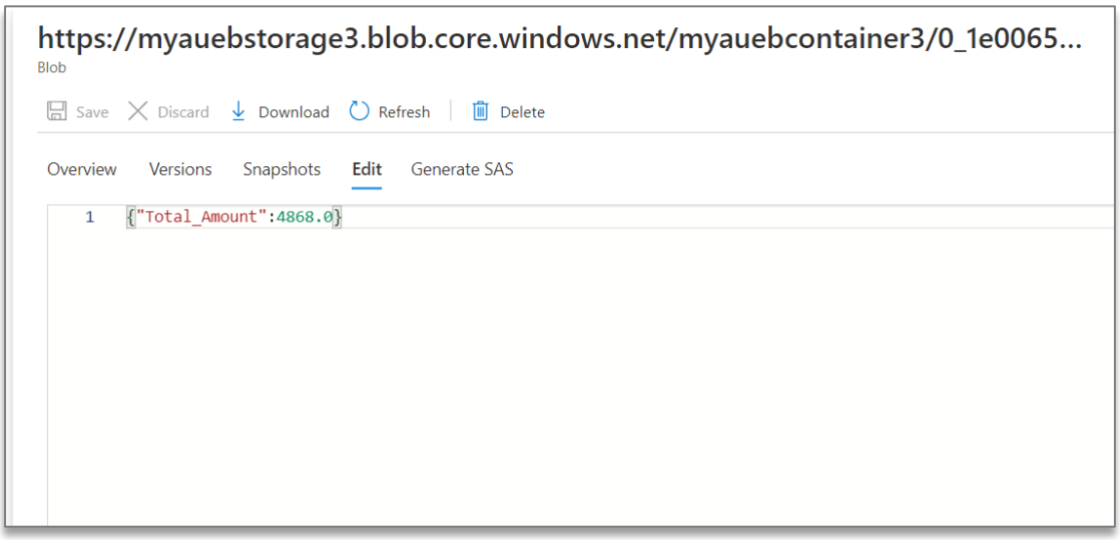

Figure 30 - Query 2 Output

Query 3: Show the total "Amount" of "Type = 1" transactions at "ATM Code = 21" of the last hour. Repeat once every 30 minutes (use a hopping window).

Query 3 – Code

```
select SUM(Amount) as Total_Amount
into [Output3Zero]
from [Input3Zero]
timestamp by EventEnqueuedUtcTime
where Type = 1 AND ATMCode = 21
group by Type, ATMCode, HoppingWindow (minute, 60, 30)
```

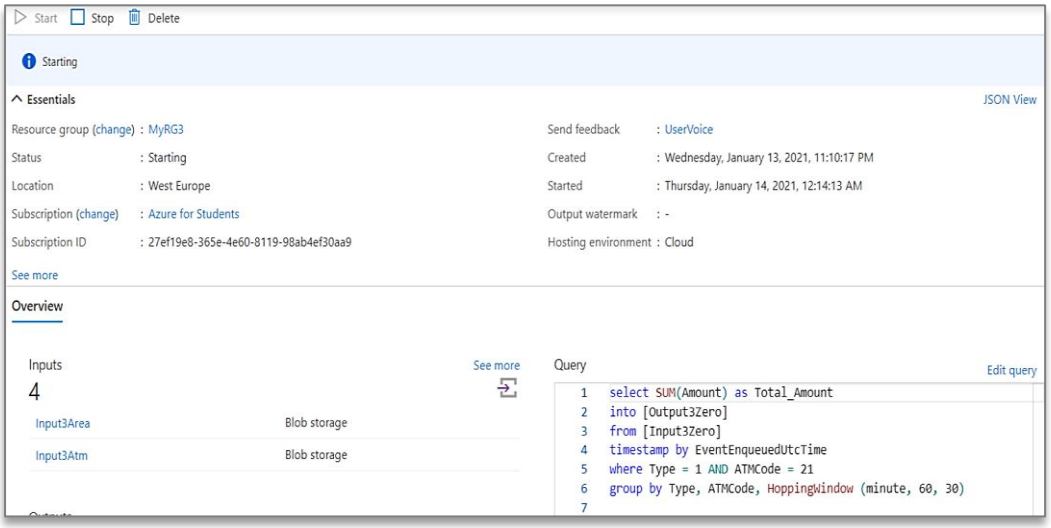

Figure 31 - Query 3 Test Results

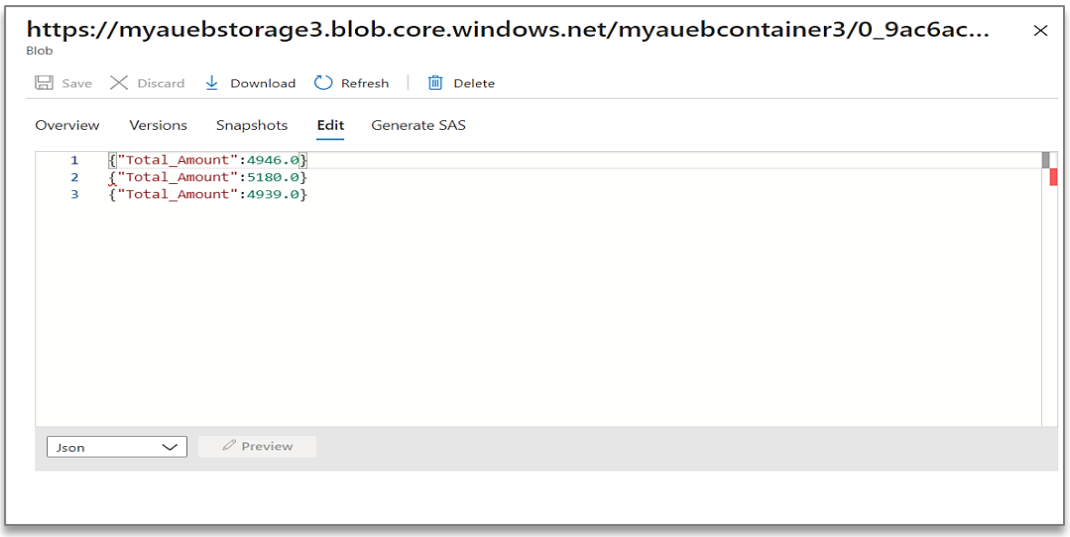

Figure 32 - Query 3 Output

Query 4: Show the total "Amount" of "Type = 1" transactions per "ATM Code" of the last one hour (use a sliding window).

Query 4 – Code

```
select SUM(Amount) as Total_Amount
into [Output3Zero]
from [Input3Zero]
timestamp by EventEnqueuedUtcTime
where Type = 1
group by Type, ATMCode, SlidingWindow (hour,1)
```

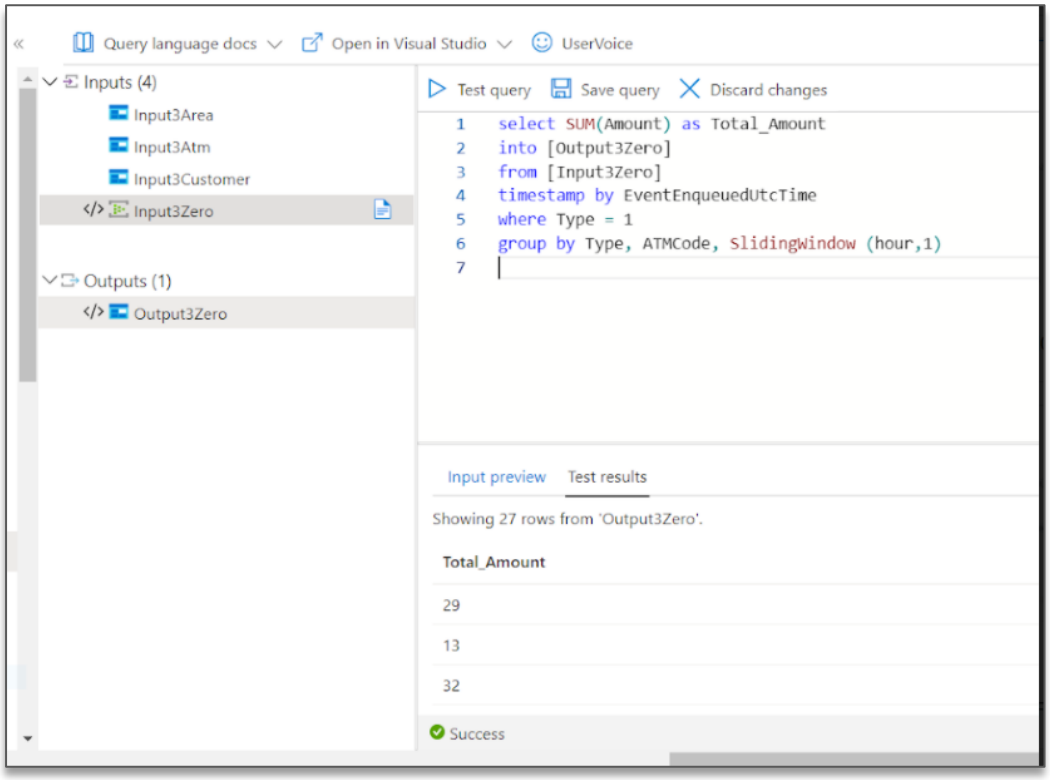

Figure 33 - Query 4 Test Results

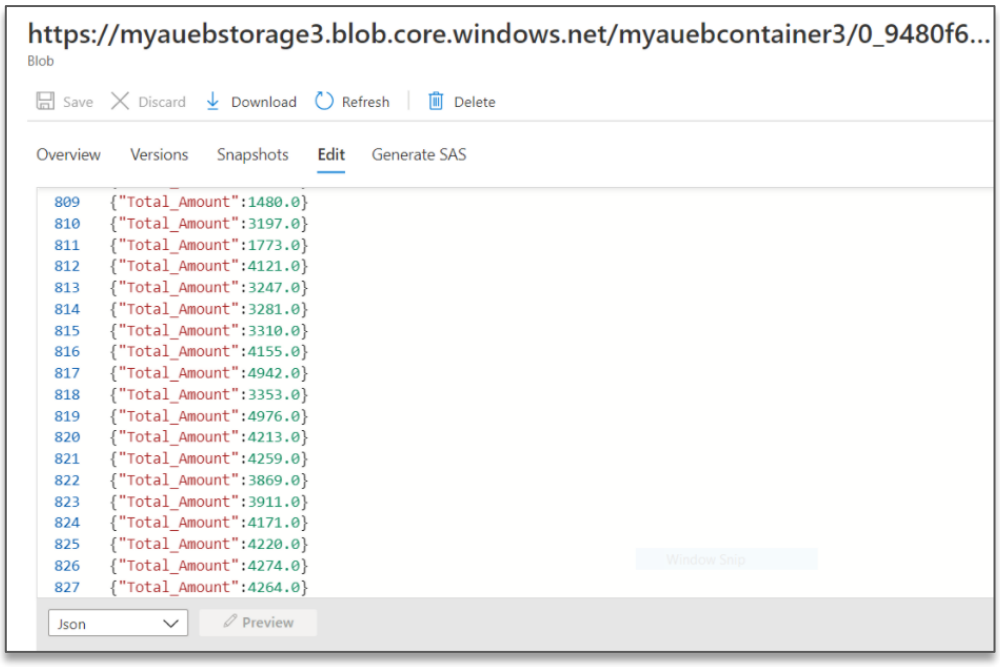

Figure 34 - Query 4 Output

Query 5: Show the total "Amount" of "Type = 1" transactions per "Area Code" of the last hour. Repeat once every hour (use a tumbling window).

Query 5 – Code

```
select Input3Zero.Type, Input3Zero.ATMCode,
SUM(Input3Zero.Amount) as Total_Amount,
into [Output3Zero]
from [Input3Zero] timestamp by EventEnqueuedUtcTime
join [Input3AtmRef]
on Input3Zero.ATMCode = Input3AtmRef.atm_code
where Input3Zero.Type = 1
group by Input3Zero.Type, Input3Zero.ATMCode,
TumblingWindow(hour, 1)
```

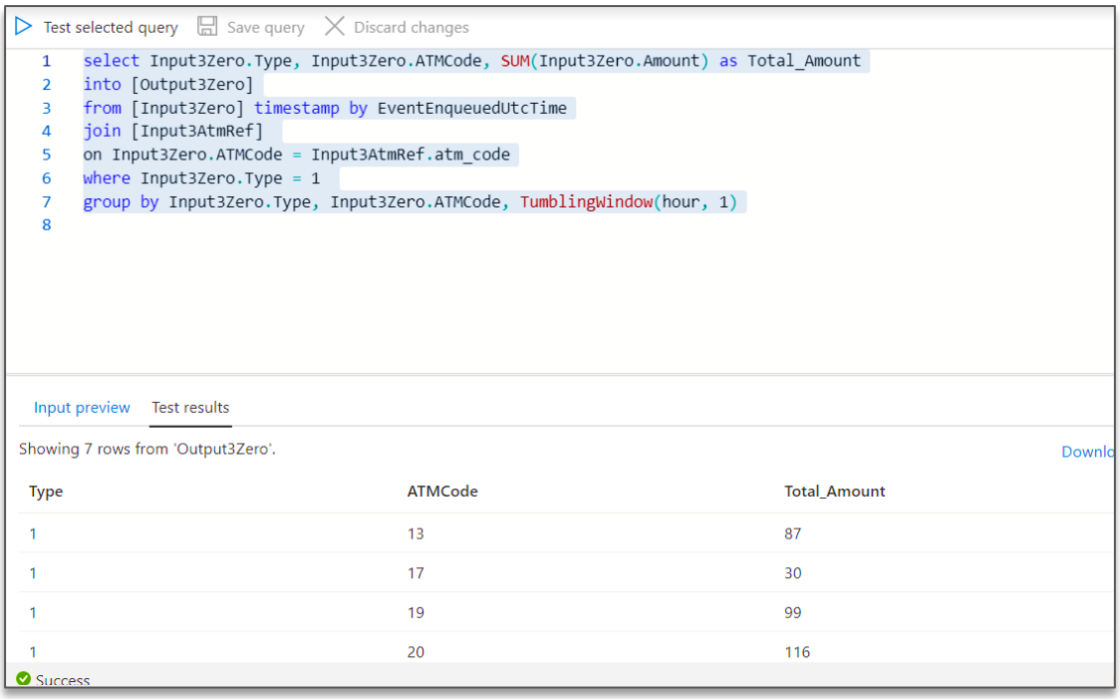

Figure 35 - Query 5 Test Results

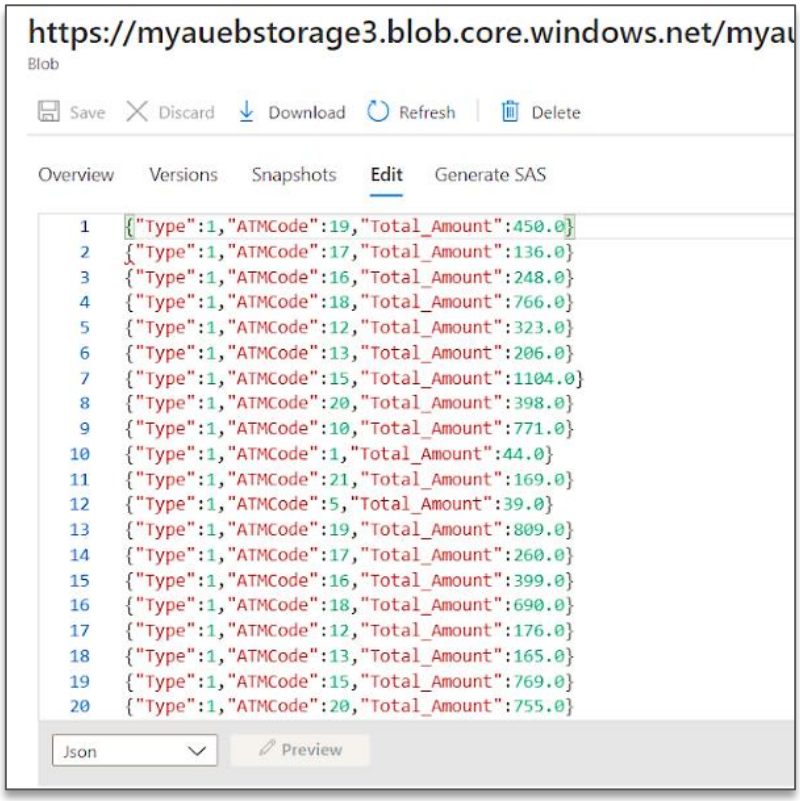

Figure 36 - Query 5 Output

Query 6: Show the total "Amount" per ATM's "City" and Customer's "Gender" of the last hour. Repeat once every hour (use a tumbling window).

Query 6 – Code

```
select Input3Zero.Type,Input3CustomerREF.gender,
SUM(Input3Zero.Amount) as Total_Amount
into [Output3Zero]
from [Input3Zero] timestamp by EventEnqueuedUtcTime
join [Input3CustomerREF]
on Input3Zero.CardNumber = Input3CustomerREF.card_number
join [Input3Arearef]
on Input3Arearef.area_code = Input3Zero.ATMCode
group by Input3Zero.Type, Input3CustomerREF.gender,
TumblingWindow(hour, 1)
```

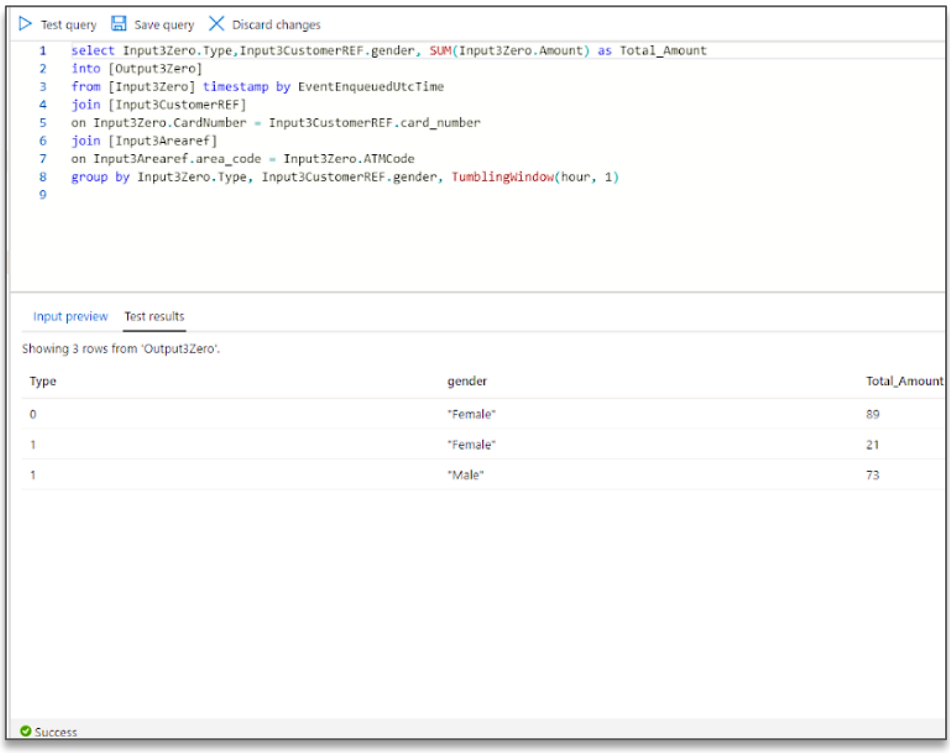

Figure 37 - Query 6 Test Results

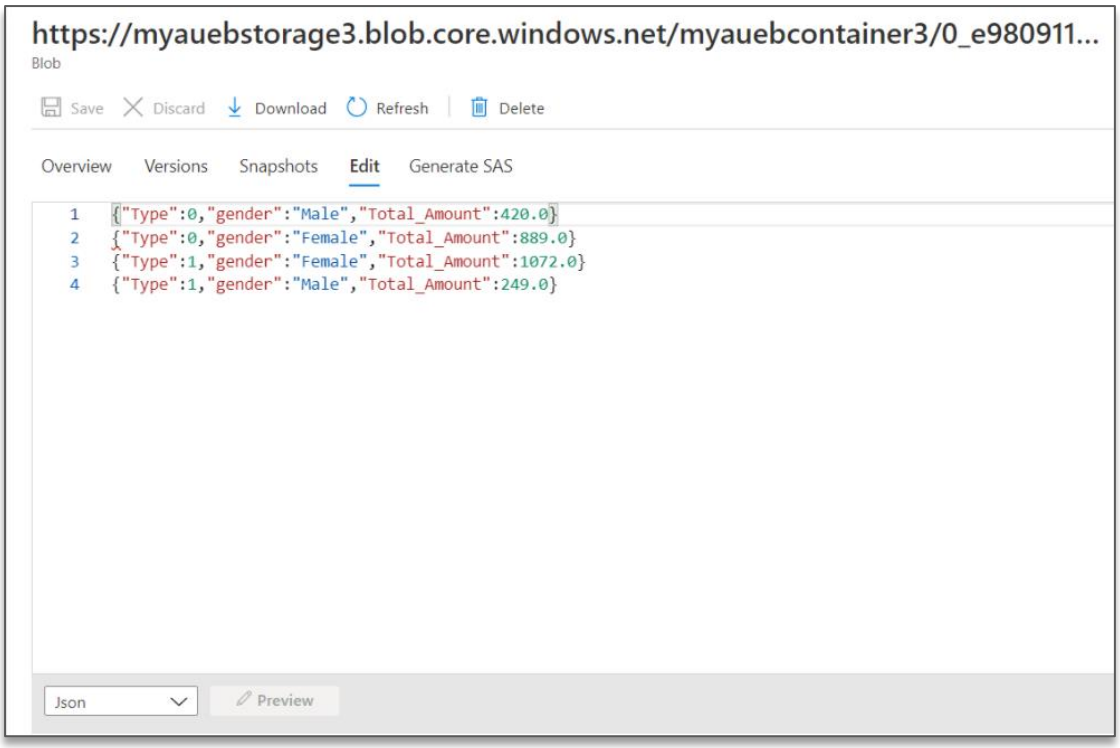

Figure 38 - Query 6 Output

Query 7: Alert (Do a simple SELECT "1") if a Customer has performed two transactions of "Type = 1" in a window of an hour (use a sliding window).

Query 7 – Code

```
select 1 as Alert, Input3CustomerREF.first_name,
Input3CustomerREF.last_name, count(Input3Zero.Type)
into [Output3Zero]
from [Input3Zero] timestamp by EventEnqueuedUtcTime
join [Input3CustomerREF]
on Input3Zero.CardNumber = Input3CustomerREF.card_number
group by Input3CustomerREF.first_name,
Input3CustomerREF.last_name,
Input3CustomerREF.card_number, Input3Zero.Type,
SlidingWindow(hour, 1)
having count(Input3Zero.Type) = 2 and Input3Zero.Type = 1
```

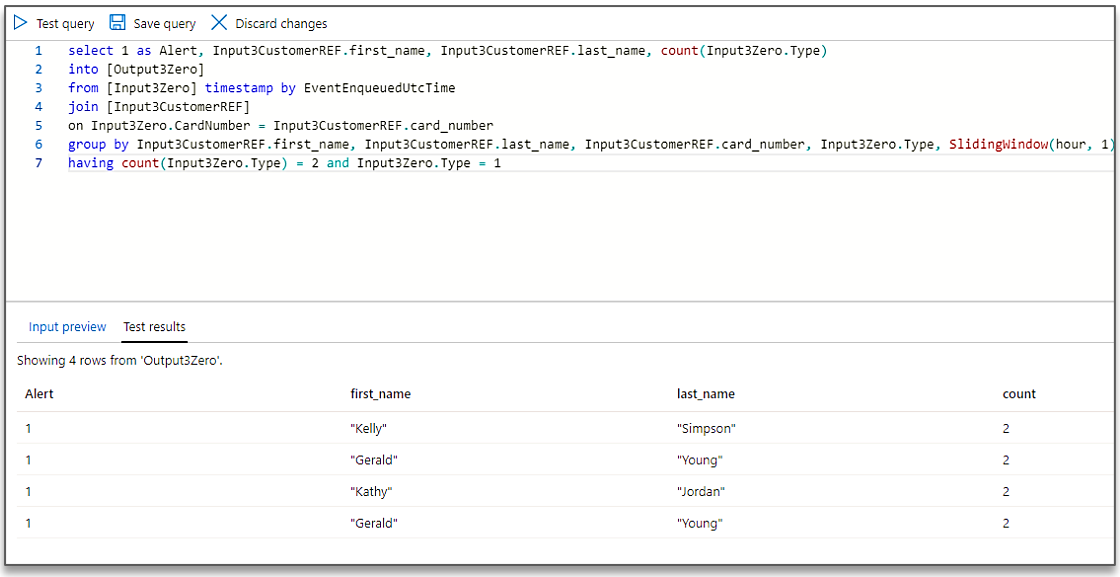

Figure 39 - Query 7 Test Results

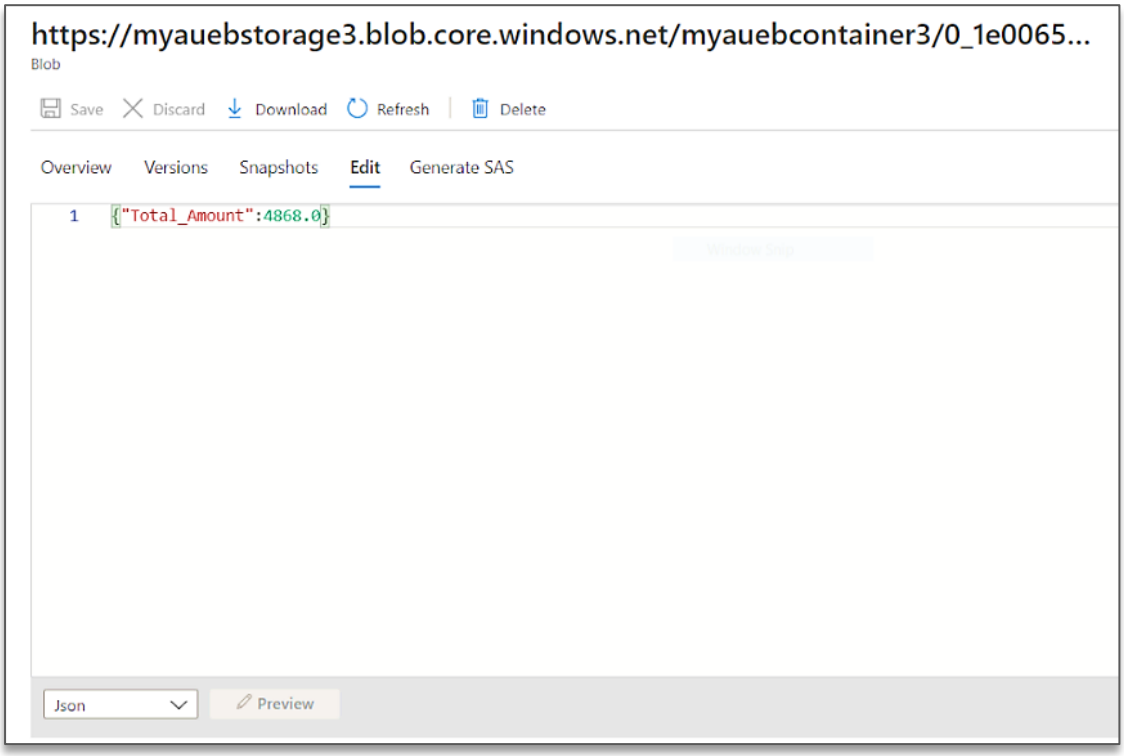

Figure 40 - Query 7 Output

Query 8: Alert (Do a simple SELECT "1") if the "Area Code" of the ATM of the transaction is not the same as the "Area Code" of the "Card Number" (Customer's Area Code) - (use a sliding window)

Query 8 – Code

```
select 1 as Alert, count(*), Input3AtmRef.area_code as
Atm_code, Input3CustomerREF.area_code as Area_code
into [Output3Zero]
from [Input3Zero] timestamp by EventEnqueuedUtcTime
join [Input3CustomerREF]
on Input3Zero.CardNumber = Input3CustomerREF.card_number
join [Input3AtmRef]
on Input3Zero.ATMcode = Input3AtmRef.atm_code
where Input3AtmRef.area_code !=
Input3CustomerREF.area_code
group by Input3AtmRef.area_code,
Input3CustomerREF.area_code, SlidingWindow (hour, 1)
```

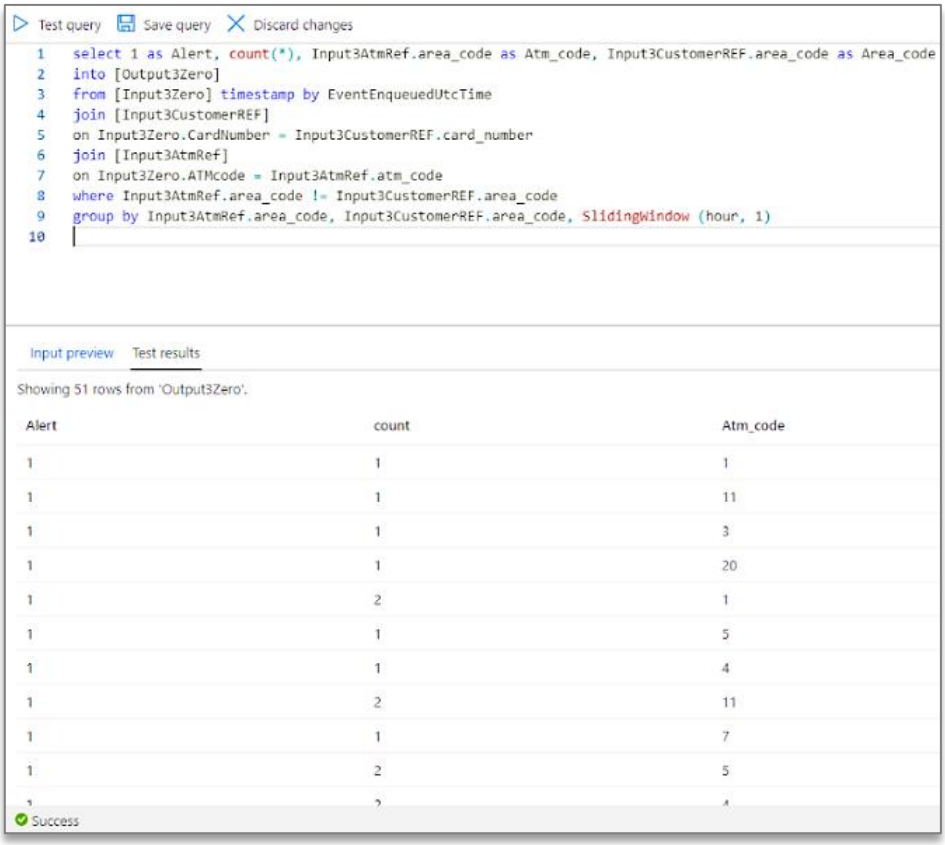

Figure 41- Query 8 Test Results

Figure 42 - Query 8 Output

References

- Azure Stream Analytics documentation
  https://docs.microsoft.com/en-us/azure/stream-analytics/